


V&VI Deep Reinforcement Learning (DRL)





Outline



1. Introduce the Deep Reinforcement Learning
2. Double Deep Q Network (Double-DQN value-based)
3. Playing Carla with Double-DQN
4. Deep Deterministic Policy Gradient (DDPG Policy Gradient)
5. Playing Carla with DDPG

1. Deep Reinforcement Learning

• Investigation:


- ① Have you heard of reinforcement learning ?
- ② Know how RL works ?
- ③ RL algorithm: DQN, Double-DQN ,Dueling DQN, Actor - Critic, DDPG ?
- ④ Reproduce the algorithm by python ?

• What is RL ?

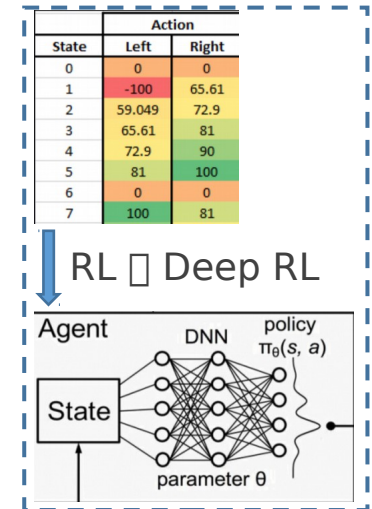
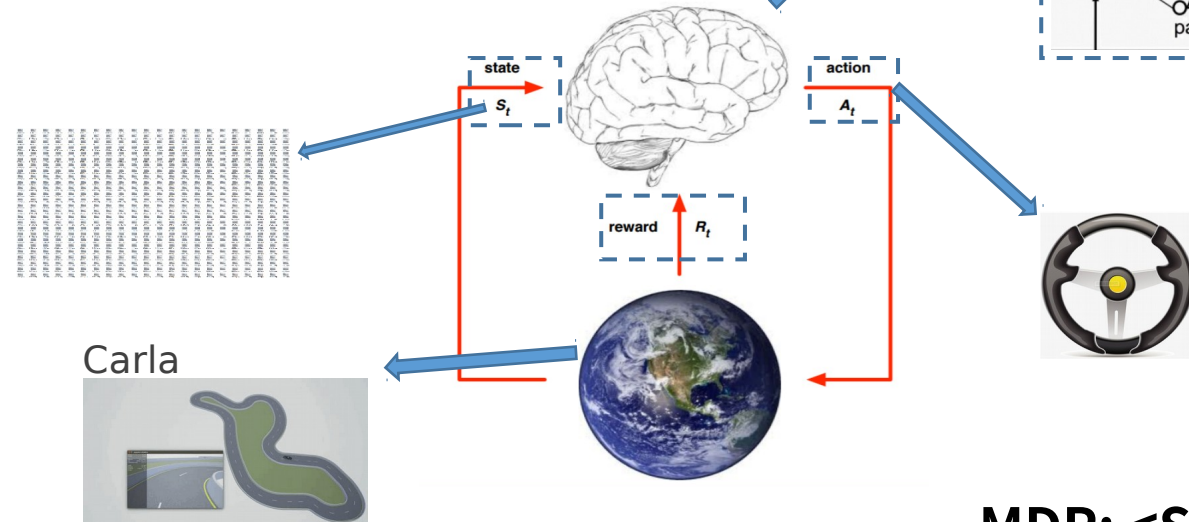
• Descript RL Process

:

Markov Decision Process



Let the robot learn **strategies** (Maximum total reward) by **interacting** with the **environment** .



MDP: $\langle S, A, P, R, \gamma \rangle$

1. Deep Reinforcement Learning

• DRL algorithms:

Value-based RL (Max the total reward)

Action discrete

~~Double DQN~~ :

$$L(\theta) = \mathbb{E}_{(s,a,r,s')} [(Q^*(s, a|\theta) - y)^2]$$

$$y = r + \gamma \max_{a'} \bar{Q}^*(s', a')$$

Improved Algorithm :

- ① Fixed target
- ② Double DQN
- ③ Dueling DQN

Instance :

Action = ['Up' , 'down' , 'left' , 'right']



Policy Gradient (Strategies)

Action Continuous

PG:

PG:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) r]$$

Basic PG algorithm:

- ① REINFORCE: Monte Carlo ;
- ② Actor-Critic: TD-error ;

Improved Algorithm :

- ① More actors A2C, A3C
- ② Replay buffer PPO **DDPG**

Instance :

Action: torque = [-2, 2]

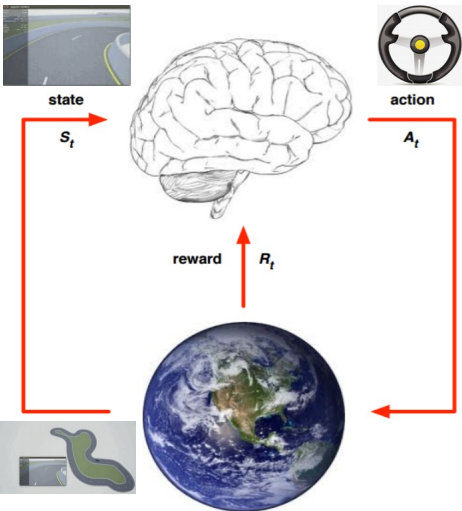


2. Value-based: DQN

• Project Analysis :


Goal : Vehicle runs **on the race road by itself**.

Two sub-project Carla environment Policy by Double DQN



• Carla Env

① Agent (Vehicle)



② State (rgb-image sensor)

③ Action (steer, brake, throttle)

• Double-DQN (Intuitively)

Algorithm 1: Double DQN Algorithm.

```

input :  $\mathcal{D}$  – empty replay buffer;  $\theta$  – initial network parameters,  $\theta^-$  – copy of  $\theta$ 
input :  $N_r$  – replay buffer maximum size;  $N_b$  – training batch size;  $N^-$  – target network replacement freq.
for episode  $e \in \{1, 2, \dots, M\}$  do
  Initialize frame sequence  $\mathbf{x} \leftarrow ()$ 
  for  $t \in \{0, 1, \dots\}$  do
    Set state  $s \leftarrow \mathbf{x}$ , sample action  $a \sim \pi_{\theta}$ 
    Sample next frame  $x'$  from environment  $\mathcal{E}$  given  $(s, a)$  and receive reward  $r$ , and append  $x^t$  to  $\mathbf{x}$ 
    if  $|\mathbf{x}| > N_r$  then delete oldest frame  $x_{t-N_r}$  from  $\mathbf{x}$  end
    Set  $s' \leftarrow \mathbf{x}$ , and add transition tuple  $(s, a, r, s')$  to  $\mathcal{D}$ , replacing the oldest tuple if  $|\mathcal{D}| \geq N_r$ 
    Sample a minibatch of  $N_b$  tuples  $(s, a, r, s') \sim \text{Unif}(\mathcal{D})$ 
    Construct target values, one for each of the  $N_b$  tuples:
    Define  $a^{\max}(s'; \theta) = \arg \max_{a'} Q(s', a'; \theta)$ 
     $y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-) & \text{otherwise.} \end{cases}$ 
    Do a gradient descent step with loss  $\|y_j - Q(s, a; \theta)\|^2$ 
    Replace target parameters  $\theta^- \leftarrow \theta$  every  $N^-$  steps
  end

```

② Y label

$y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-) & \text{otherwise.} \end{cases}$

Do a gradient descent step with loss $\|y_j - Q(s, a; \theta)\|^2$

Replace target parameters $\theta^- \leftarrow \theta$ every N^- steps

end

end

② Y label

$$Q(s, a; \theta) = r + \gamma Q(s', \arg \max_{a'} Q(s', a'; \theta); \theta^-)$$

Reward
Next state Q-value

Double

2. Value-based: DQN

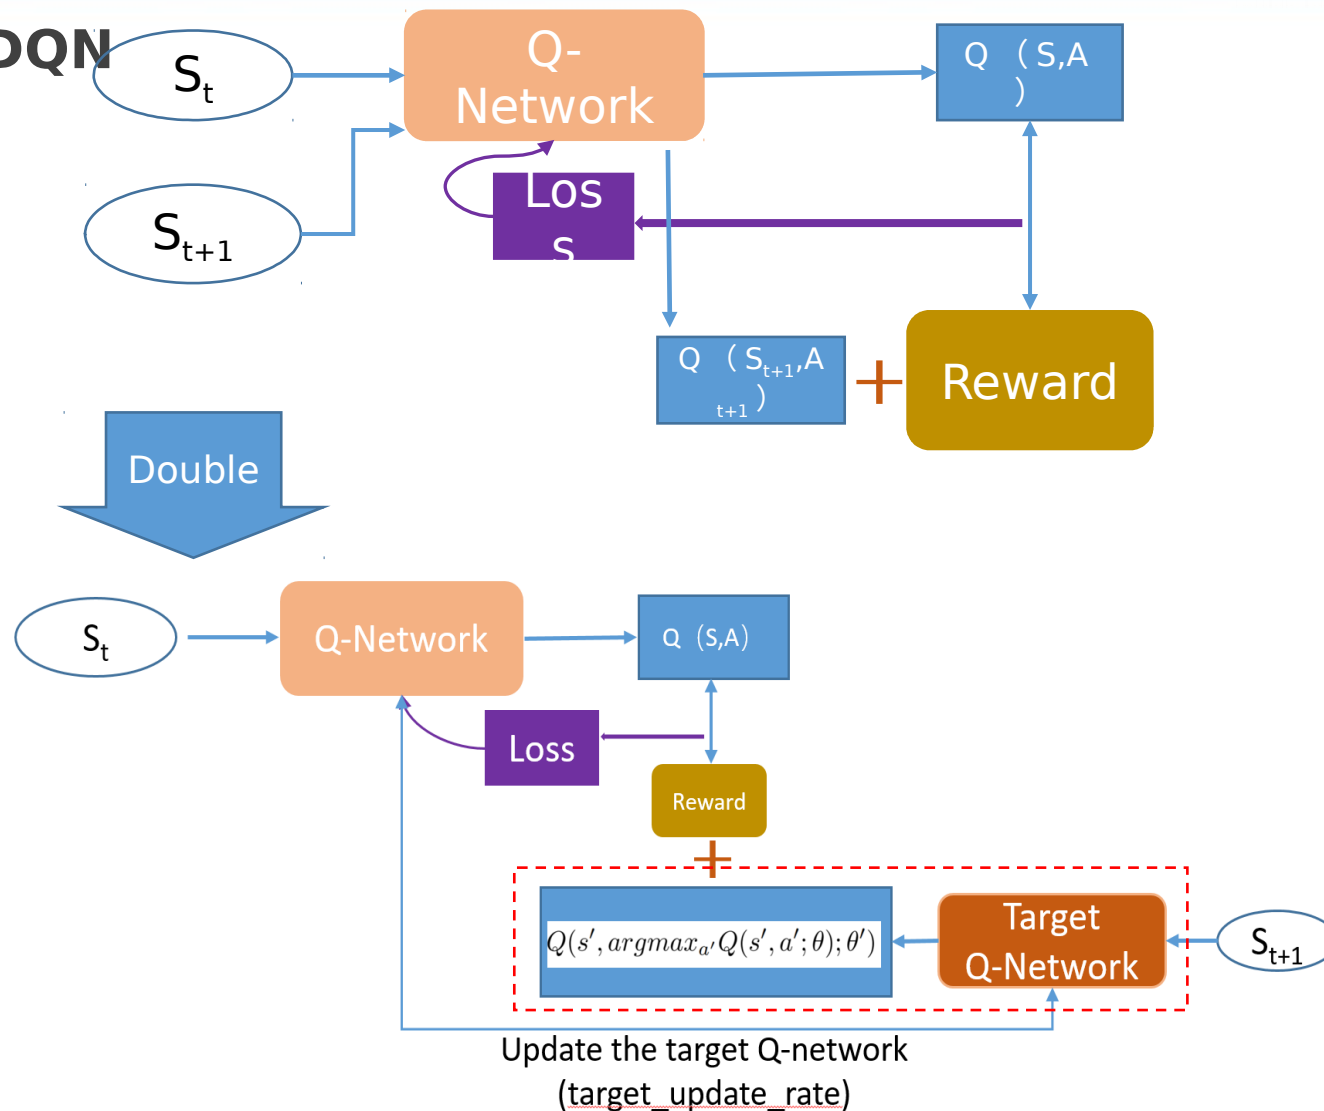
• Naïve DQN - Fixed DQN - Double DQN

① Naïve DQN ② Fixed DQN

$$Q(s, a; \theta) = r + \gamma Q(s', \argmax_{a'} Q(s', a'; \theta); \theta)$$

③ Double DQN

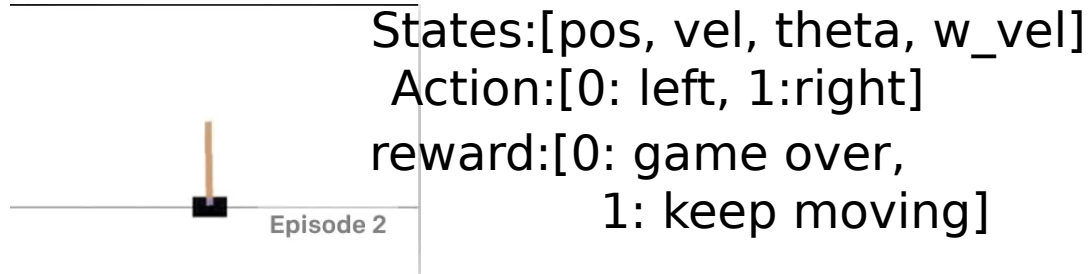
$$Q(s, a; \theta) = r + \gamma Q(s', \argmax_{a'} Q(s', a'; \theta); \theta')$$





2. Value-based: DQN

• Double DQN: CartPole



Algorithm 1: Double DQN Algorithm.

```


input :  $\mathcal{D}$  – empty replay buffer;  $\theta$  – initial network parameters,  $\theta^-$  – copy of  $\theta$ 
input :  $N_r$  – replay buffer maximum size;  $N_b$  – training batch size;  $N^-$  – target network replacement freq.
for episode  $e \in \{1, 2, \dots, M\}$  do
  Initialize frame sequence  $\mathbf{x} \leftarrow ()$ 
  for  $t \in \{0, 1, \dots\}$  do
    Set state  $s \leftarrow \mathbf{x}$ , sample action  $a \sim \pi_{\theta}$ 
    Sample next frame  $x^t$  from environment  $\mathcal{E}$  given  $(s, a)$  and receive reward  $r$ , and append  $x^t$  to  $\mathbf{x}$ 
    if  $|\mathbf{x}| > N_r$  then delete oldest frame  $x_{t_{min}}$  from  $\mathbf{x}$  end
    Set  $s' \leftarrow x^t$ , and add transition tuple  $(s, a, r, s')$  to  $\mathcal{D}$ ,
      replacing the oldest tuple if  $|\mathcal{D}| \geq N_r$ 
    Sample a minibatch of  $N_b$  tuples  $(s, a, r, s') \sim \text{Unif}(\mathcal{D})$ 
    Construct target values, one for each of the  $N_b$  tuples:
    Define  $a^{\max}(s'; \theta) = \arg \max_{a'} Q(s', a'; \theta)$ 
    
$$y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-) & \text{otherwise.} \end{cases}$$

    Do a gradient descent step with loss  $\|y_j - Q(s, a; \theta)\|^2$ 
    Replace target parameters  $\theta^- \leftarrow \theta$  every  $N^-$  steps
  end
end
  
```

① Replay buffer

transition(s, a, r, s'), memory_size, mini-batch

② Epsilon Greedy (Exploration-- Exploitation)



Exploration

Random(0,1) < red point value

Exploitation

Random(0,1) > red point value

③ Interact with environment


Action □ Next state □ reward □ Done

④ Double NETWORK

Q-network (Fully connected) predicts value & target value

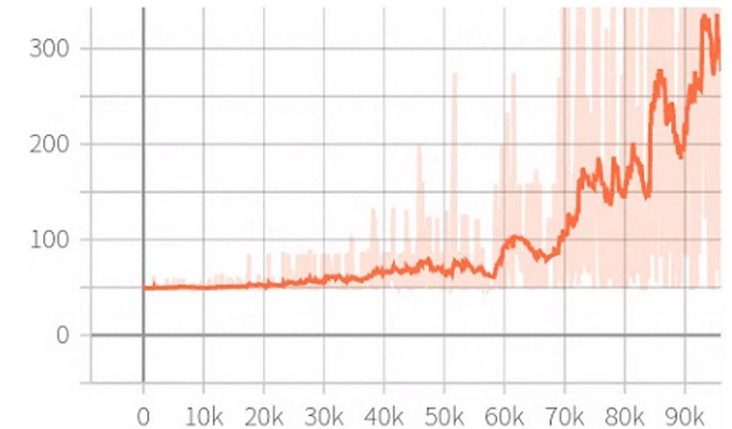
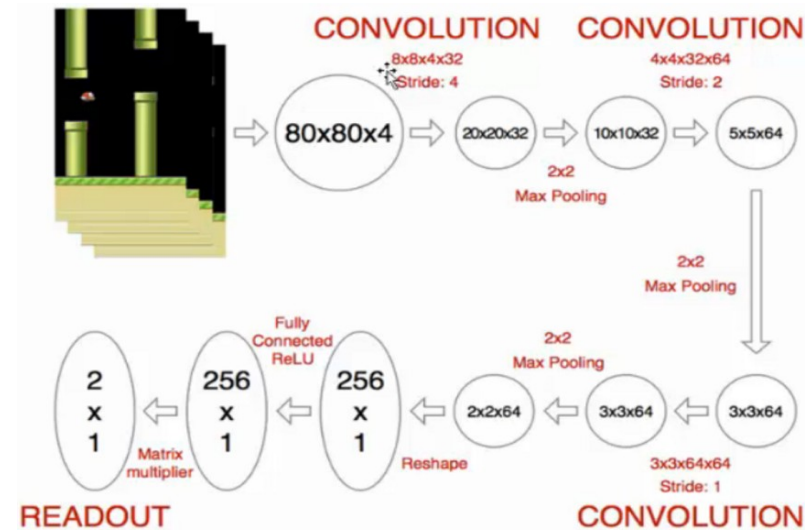
2. Value-based: DQN

- Double DQN: Flappy bird (image)



States ? Action ? Reward ? Done ?

DQN: CNN



Episode reward

Tricks:

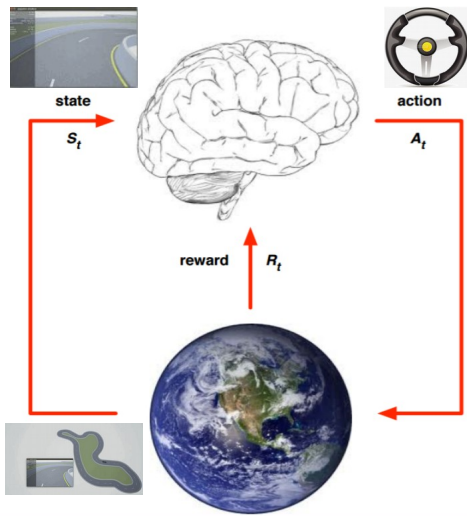
- ① Rgb image → gray image;
- ② Resize image;
- ③ 4 frame;

HyperParameter:

- ① Max episode;
- ② Memory size;
- ③ Batch size;
- ④ Epsilon decay;
- ⑤ Target network update
- ⑥ Learning rate;

2. Value-based: DQN

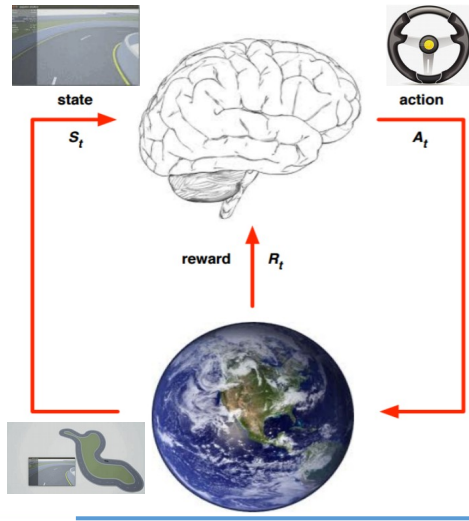
- Carla step by step :



- ① Create a Carla Environment;
 - a) Spawn a vehicle
 - b) Spawn a rgb-camera attaching to vehicle
- ② Build a CNN network (input: image);
 - a) CNN network
 - b) State dimension and action dimension
- ③ Achieve Double DQN algorithm to train CNN network;
 - a) Epsilon-greedy policy - action;
 - b) Interact with env
 - c) Transition (s, a, r, s')
 - d) Training network by replay buffer
- ④ Use the CNN (Policy) to control the vehicle;
 - a) Load the CNN params;
 - b) Interact with Carla env;

2. Value-based: DQN

• Create a Carla Environment :



CarlaEnv Class:

- ① Reset : spawn vehicle and sensor for each episode ;
- ② Collision_data: Vehicle have a collision and this episode is done;
- ③ Process_image: Record a env state;
- ④ Step: agent interacts with the Carla simulator;
 - a) Apply a control based on the action;
 - b) Design a reward for the action;
 - c) The episode is done or not;
- ⑤ Find the startpoint;
- ⑥ Destroy the actors (vehicle and sensor);

```
class CarEnv:

    SHOW_CAM = SHOW_PREVIEW
    STEER_AMT = 0.3
    im_width = IMG_WIDTH
    im_height = IMG_HEIGHT
    actor_list = []
    front_camera = None
    collision_hist = []

    def __init__(self):
        self.client = carla.Client('localhost', 2000)
        self.client.set_timeout(3.0)
        self.world = self.client.get_world()
        self.server_clock = pygame.time.Clock()
        # the start way_point
        self.start_point = self.set_start_waypoint()
        # states
        self.states = None
        self.surface = None

    def set_start_waypoint(self):...

    def reset(self):...

    def collision_data(self, event):...

    def process_img(self, image):...

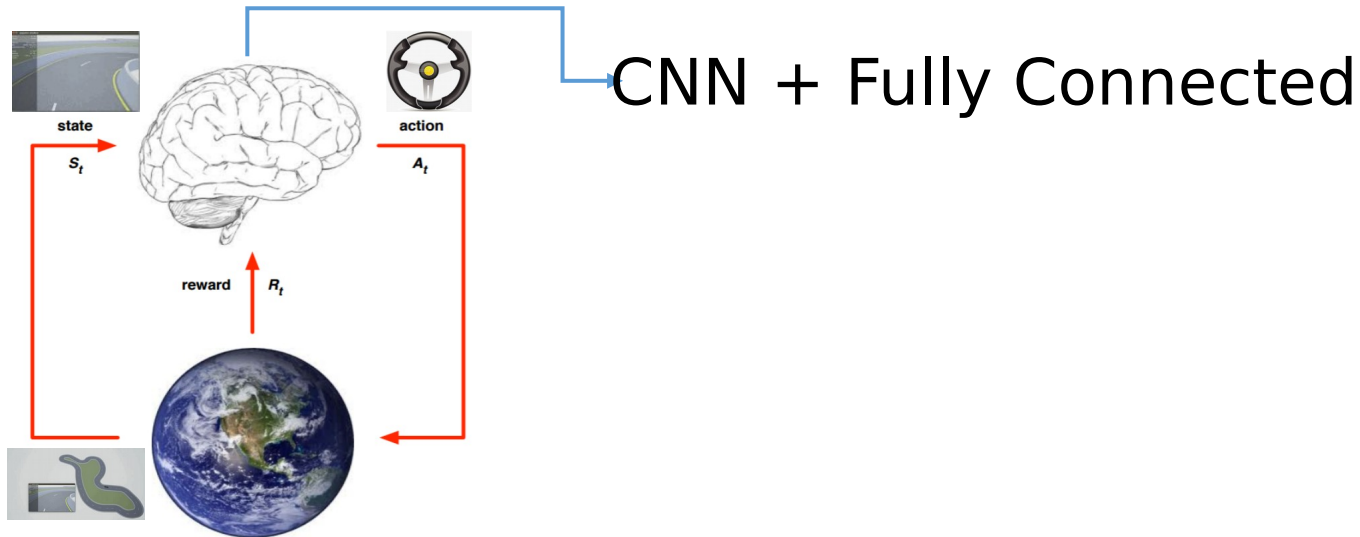
    def step(self, action):...

    def running_demo(self):...

    def destroy_actors(self):...
```

2. Value-based: DQN

- Build a CNN network :



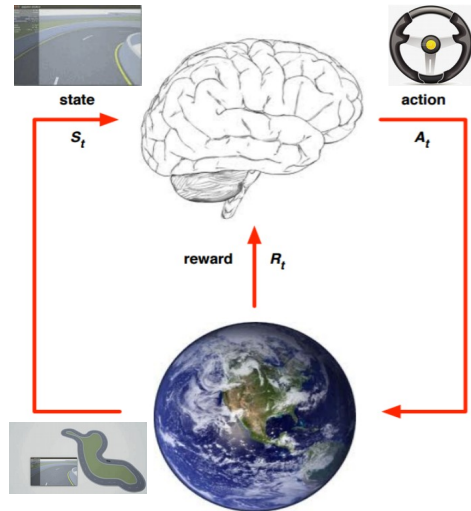
```
class Network(nn.Module):

    def __init__(self, image_channel=1, output_dim=3):
        super(Network, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=image_channel, out_channels=24, kernel_size=5, stride=(2, 2))
        self.conv1_bn = nn.BatchNorm2d(24)
        self.conv2 = nn.Conv2d(in_channels=24, out_channels=36, kernel_size=5, stride=(2, 2))
        self.conv2_bn = nn.BatchNorm2d(36)
        self.conv3 = nn.Conv2d(in_channels=36, out_channels=48, kernel_size=5, stride=(2, 2))
        self.conv3_bn = nn.BatchNorm2d(48)
        self.conv4 = nn.Conv2d(in_channels=48, out_channels=64, kernel_size=3, stride=(1, 1))
        self.conv4_bn = nn.BatchNorm2d(64)
        self.conv5 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, stride=(1, 1))
        self.fc1 = nn.Linear(in_features=1280, out_features=256)
        self.fc2 = nn.Linear(in_features=256, out_features=10)
        self.Adan = nn.Linear(10, output_dim)
        self.V = nn.Linear(10, 1)

    def forward(self, x):
        # reshape size
        x = F.interpolate(x, size=[60, 120], mode="bilinear", align_corners=False)
        # conv 1
        x = self.conv1_bn(F.elu(self.conv1(x)))
        # conv 2
        x = self.conv2_bn(F.elu(self.conv2(x)))
        # conv 3
        x = self.conv3_bn(F.elu(self.conv3(x)))
        # conv 4
        x = self.conv4_bn(F.elu(self.conv4(x)))
        # Flatten batch * dim
        x = x.view(-1, 1280)
        # fc 1
        x = F.dropout(F.relu(self.fc1(x)), 0.2)
        # fc 2
        x = F.dropout(F.relu(self.fc2(x)), 0.5)
        # output
        adv = self.Adan(x)
        v = self.V(x)
        adv_average = torch.mean(adv, dim=-2, keepdim=True)
        return v + (adv-adv_average)
```

2. Value-based: DQN

• Training Network by Double DQN:



Obtain the current state (rgb-image)

Epsilon-greedy

Agent interacts with Carla Env

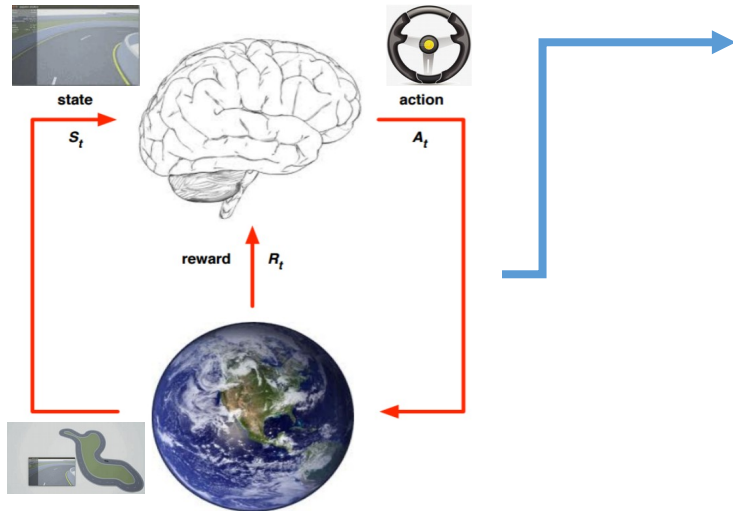
Replay Buffer

Collect data

```
# Iterate over episodes
for episode_i in range(1, EPISODES + 1):
    # Reset environment and get initial state
    current_state = env.reset()
    # Reset flag and start iterating until episode ends
    done = False
    # Play for given number of seconds only
    while True:
        time_step += 1
        # This part stays mostly the same, the change is to query a model for Q values
        if np.random.random() > epsilon:
            # Get action from Q table
            action = agent.get_qs(current_state).argmax(dim=-1).detach().to('cpu').numpy()[0]
        else:
            # Get random action [0, 1, 2] -> steer: left, middle, right
            action = np.random.randint(0, 3)
            # This takes no time, so we add a delay matching 60 FPS (prediction above takes longer)
            # time.sleep(1 / FPS)
        # interact with the carla Env
        new_state, reward, done, _ = env.step(action)
        if time_step > 20:
            # waiting for spawning the vehicle
            # Every step we update replay memory
            agent.update_replay_memory((current_state, action, reward, new_state, done))
        # transform the state
        current_state = new_state
        episode_reward += reward
    if done:
        break
```

2. Value-based: DQN

- Training Network by Double DQN:



Training Model

```
# -----#
# ---- Initialization the training model ---- #
# -----#
# Start training thread and wait for training to be initialized
trainer_thread = Thread(target=agent.train_in_loop, daemon=True)
trainer_thread.start()
```

Multi
Threads

Q-value(s,
a)

Target Q-value

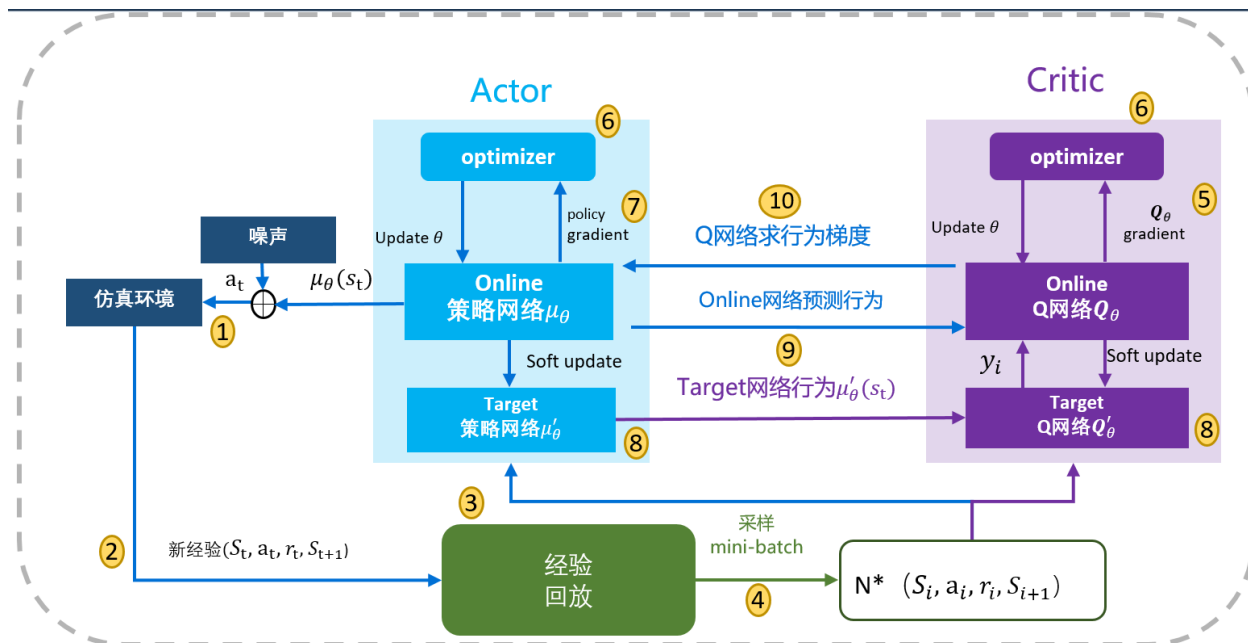
Loss
function

Pytorch
Gradient
Descent

```
# --- double dqn --- #
# 1. current q value
q_value = self.q_model(state_b)
max_q_value = q_value.gather(1, action_b)
# --- dqn target --- #
next_q_action = self.q_model(next_state_b).argmax(dim=-1).unsqueeze(dim=-1)
target_q = self.target_q_model(next_state_b).gather(1, next_q_action)
target_q_value = reward_b + DISCOUNT * target_q * done_b
# loss - gradient - update
loss = F.smooth_l1_loss(max_q_value, target_q_value)
# grad zero
self.optimizer.zero_grad()
# cal gradient
loss.backward()
# # clip
nn.utils.clip_grad_norm_(self.q_model.parameters(), max_norm=0.5)
# update
self.optimizer.step()
```


Deep Deterministic Policy Gradient (DDPG Policy Gradient)

DDPG



- ① Action;
- ② Environment interaction new transition;
- ③ Replay buffer;
- ④ Mini-batch;
- ⑤ Critic DQN gradient descent;
- ⑥ Adam optimizer;
- ⑦ Actor gradient ascent;
- ⑧ Update the target network;
- ⑨ Critic input: actor predicts action;
- ⑩ Critic update actor parameter; θ^μ ;

Critic DQN

MSE :

$$L(\theta^Q) = \frac{1}{N} \sum_{i=1}^N (y_i - Q(s_i, a_i | \theta^Q))^2$$

: target DQN label

$$y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1} | \theta^{\mu'})) - Q(s_i, a_i | \theta^Q)$$

Actor Policy Network

$$\nabla_{\theta^\mu} J = \frac{1}{N} \sum_i [\nabla_{\theta^\mu} \mu(s_i | \theta^\mu) * \nabla_a Q(s, a | \theta^Q) |_{a=\mu(s), s=s_i}]$$

$$L(\theta^\mu) = -\frac{1}{N} \sum_i Q(s, a) |_{a=\mu(s)}$$

Goal



Control Method

- ① Keyboard;
- ② PID;
- ③ Behavior Cloning;
- ④ Reinforcement learning (DQN);

Sensor

- ① RGB-image;
- ② Depth-image;
- ③ Lidar;

Just the beginning

- Improve the RL control
 - ✓ Steer, Throttle, brake and more ;
 - ✓ Policy gradient ;
- Autonomous driving license
 - ✓ Side parking ;
 - ✓ Reversing into the garage ;
 - ✓ Right angle bend;
- Multi-agent
 - ✓ V2V ;
 - ✓ V2X ;

